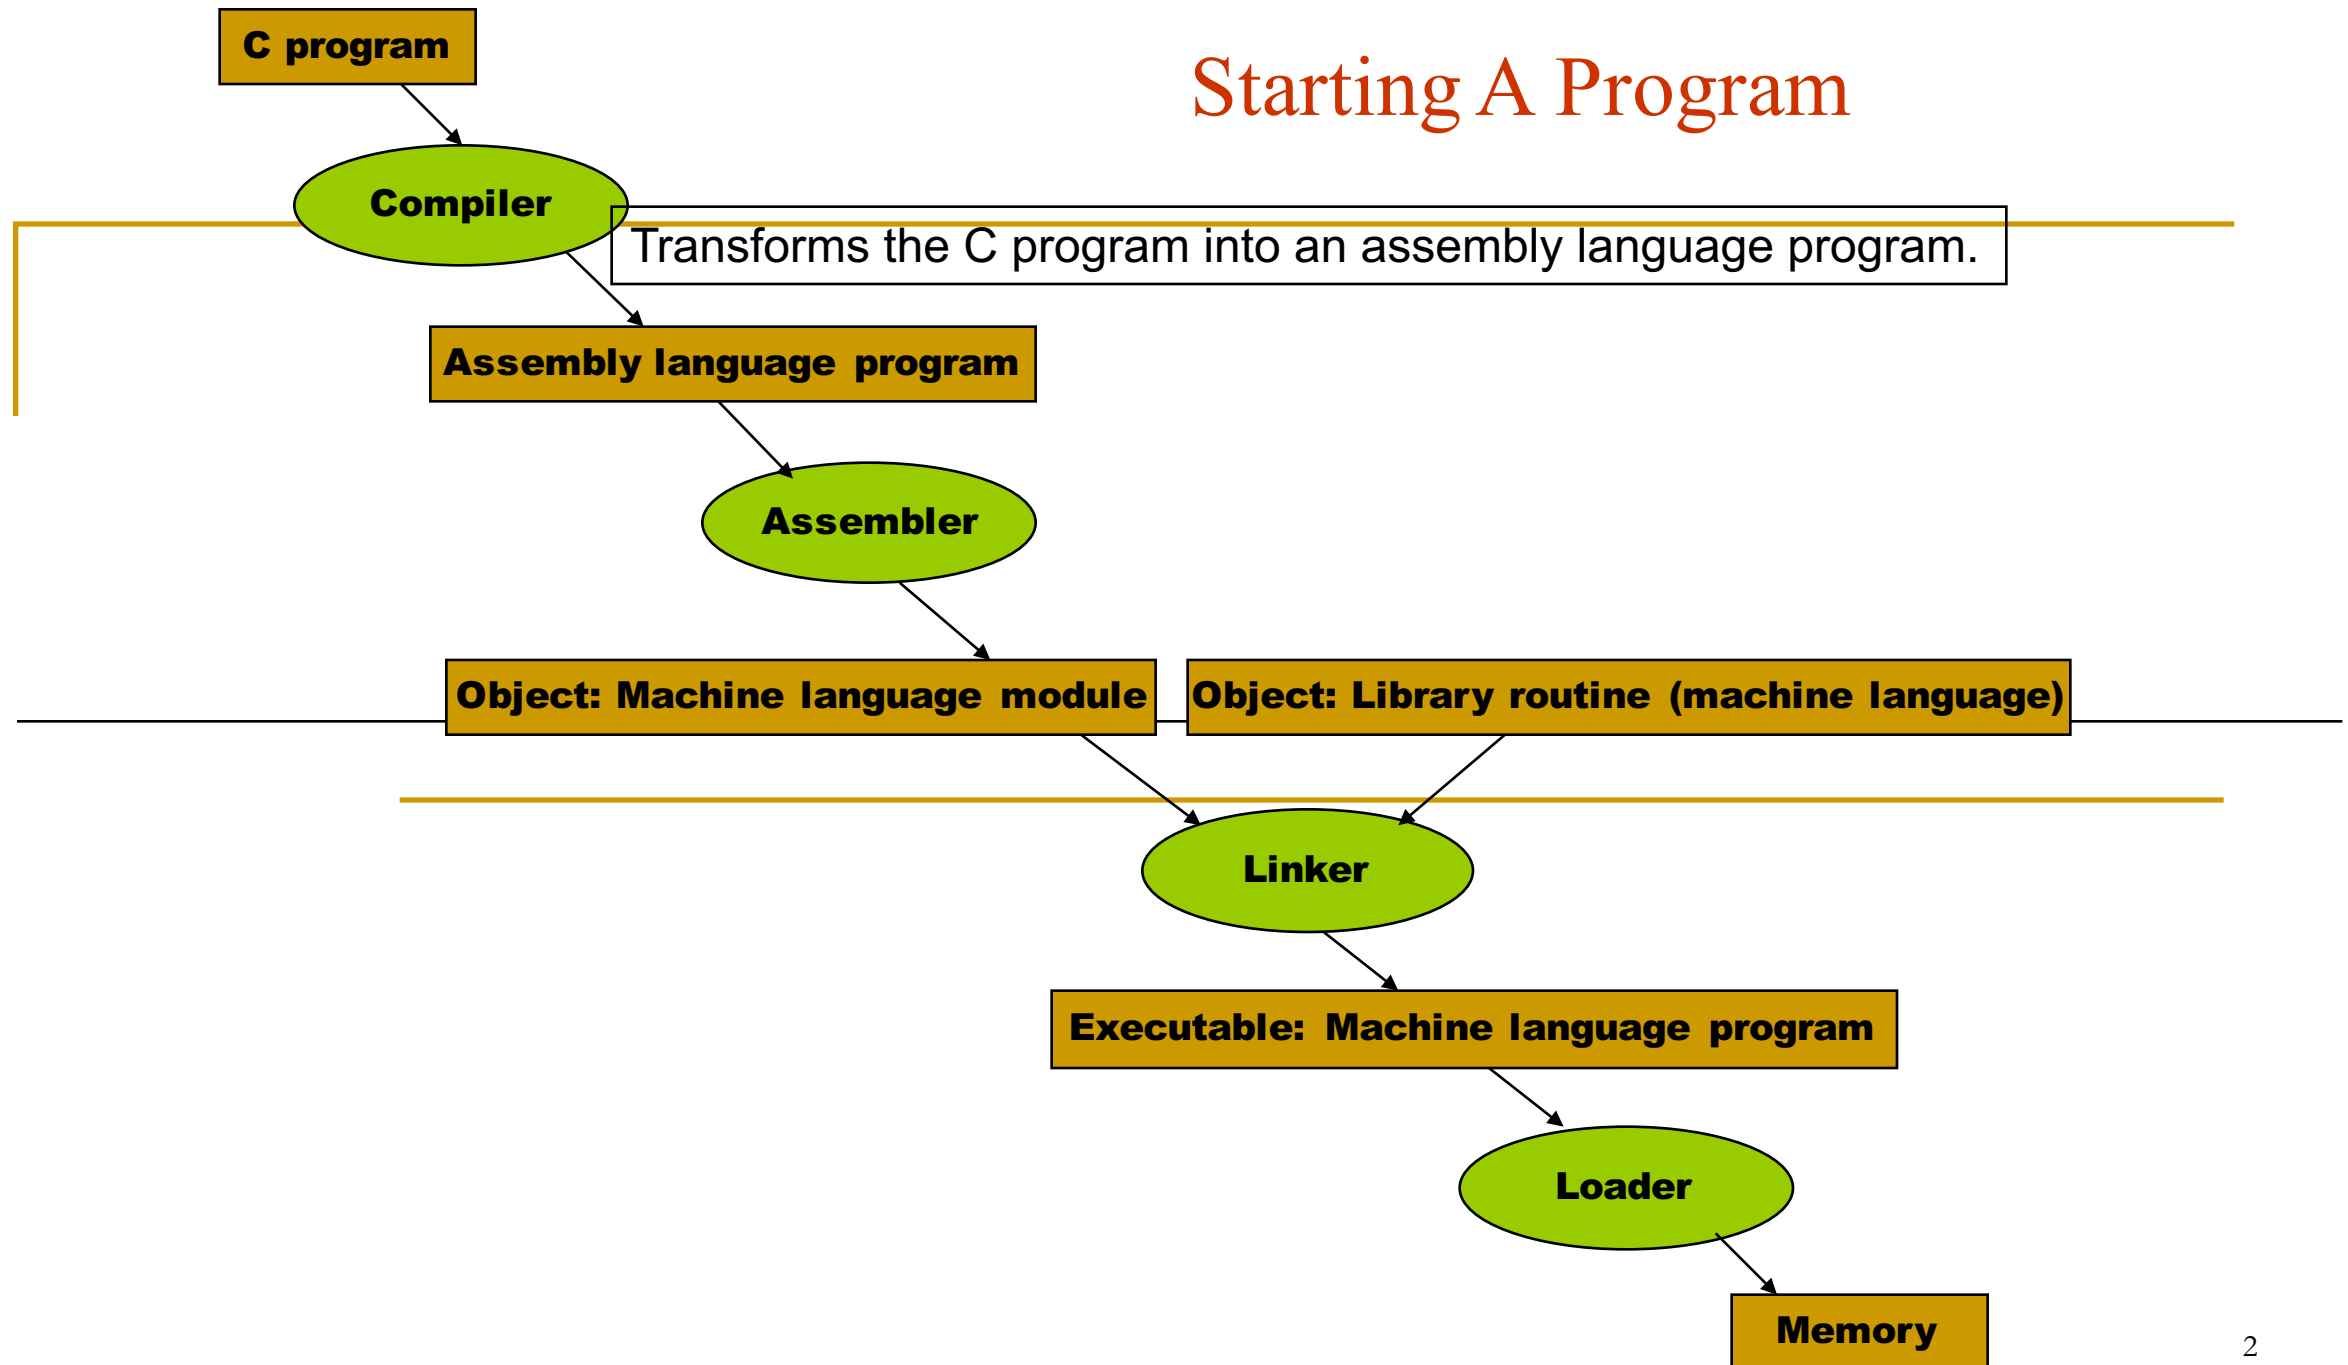


# Linker & Loader

# Starting A Program



# Assembler

---

## ■ Assembler

- ❑ The assembler turns the assembly language program into an object file.
- ❑ Symbol table: A table that matches names of labels to the addresses of the memory words that instruction occupy.

Object file header			
	Name	Procedure A	
	Text size	100 <sub>hex</sub>	
	Data size	20 <sub>hex</sub>	
Text segment	Address	Instruction	
	0	ld x10,0(x3)	
	4	jal x1, 0	
	...	...	
Data segment	0	(X)	
	...	...	
Relocation information	Address	Instruction Type	Dependency
	0	ld	X
	4	jal	B
Symbol table	Label	Address	
	X	—	
	B	—	

ld x10, x  
jal x1, B

# Assembler (cont.)

---

- Pseudoinstruction: a common variation of assembly language instructions often treated as if it were an instruction in its own right.
  - `li x9, 123` -> `addi x9, x0, 123`
  - `mv x10, x11` -> `addi x10, x11, 0`

# Linker (Link editor)

---

- Linker takes all the independently assembled machine language programs and “stitches” them together to produce an executable file that can be run on a computer.
- There are three steps for the linker:
  1. Place code and data modules symbolically in memory.
  2. Determine the addresses of data and instruction labels.
  3. Patch both the internal and external references.

Object file header			
	Name	Procedure A	
	Text size	100hex	
	Data size	20hex	
Text segment	Address	Instruction	
	0	ld x10, 0(x3)	
	4	jal x1, 0	
	...	...	
Data segment	0	(X)	
	...	...	
Relocation information	Address	Instruction Type	Dependency
	0	ld	X
	4	jal	B
Symbol table	Label	Address	
	X	—	
	B	—	


Object file header			
	Name	Procedure B	
	Text size	200hex	
	Data size	30hex	
Text segment	Address	Instruction	
	0	sd \$x11, 0(x3)	
	4	jal x1, 0	
	...	...	
Data segment	0	(Y)	
	...	...	
Relocation information	Address	Instruction Type	Dependency
	0	sd	Y
	4	jal	A
Symbol table	Label	Address	
	Y	—	
	A	—	

SP → 0000 003f ffff fff0<sub>hex</sub>

x3 → 0000 0000 1000 0000<sub>hex</sub>

PC → 0000 0000 0040 0000<sub>hex</sub>

0



Executable file header		
	Text size	300hex
	Data size	50hex
Text segment	Address	Instruction
	0040 0000hex	ld \$x10, 0 (x3)
	0040 0004hex	Jal x1, 252 <sub>ten</sub>
	...	...
	0040 0100hex	sd \$x11, 32 (x3)
	0040 0104hex	jal x1, -260 <sub>ten</sub>
	...	...
Data segment	Address	
	1000 0000hex	(X)
	...	...
	1000 0020hex	(Y)
	...	...

# Loader

- Read the executables file header to determine the size of the text and data segments
- Creates an address space large enough for the text and data
- Copies the instructions and data from the executable file into memory
- Copies the parameters (if any) to the main program onto the stack
- Initializes the machine registers and sets the stack pointer the first free location
- Jump to a start-up routine

```
main();
```

```
_start_up:
```

```
lw  a0, offset($sp)  ## load arguments
```

```
jal  main;
```

```
exit
```



# Loading a Program

- Load from image file on disk into memory

1. Read header to determine segment sizes

2. Create virtual address space

3. Copy text and initialized data into memory

- Or set page table entries so they can be faulted in

4. Set up arguments on stack

5. Initialize registers (including sp, fp, gp)

6. Jump to startup routine

- Copies arguments to x10, ... and calls main
- When main returns, do exit syscall

```
main();
```

```
_start_up:
```

```
lw x10, offset($sp) ## load arguments
```

```
jal x1, main;
```

```
exit
```

# Dynamically Linked Libraries (DLL)

---

- Disadvantages with traditional statically linked library
  - ❑ Library updates
  - ❑ Loading the whole library even if all of the library is not used
- Dynamically linked library
  - ❑ The libraries are not linked and loaded until the program is run.
  - ❑ Lazy procedure linkage
    - Each routine is linked only after it is called.

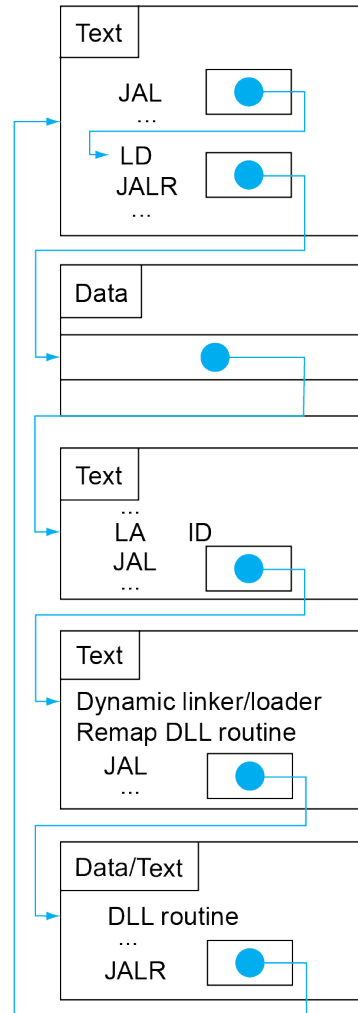
# Lazy Linkage

Indirection table


Stub: Loads routine ID,  
Jump to linker/loader

Linker/loader code

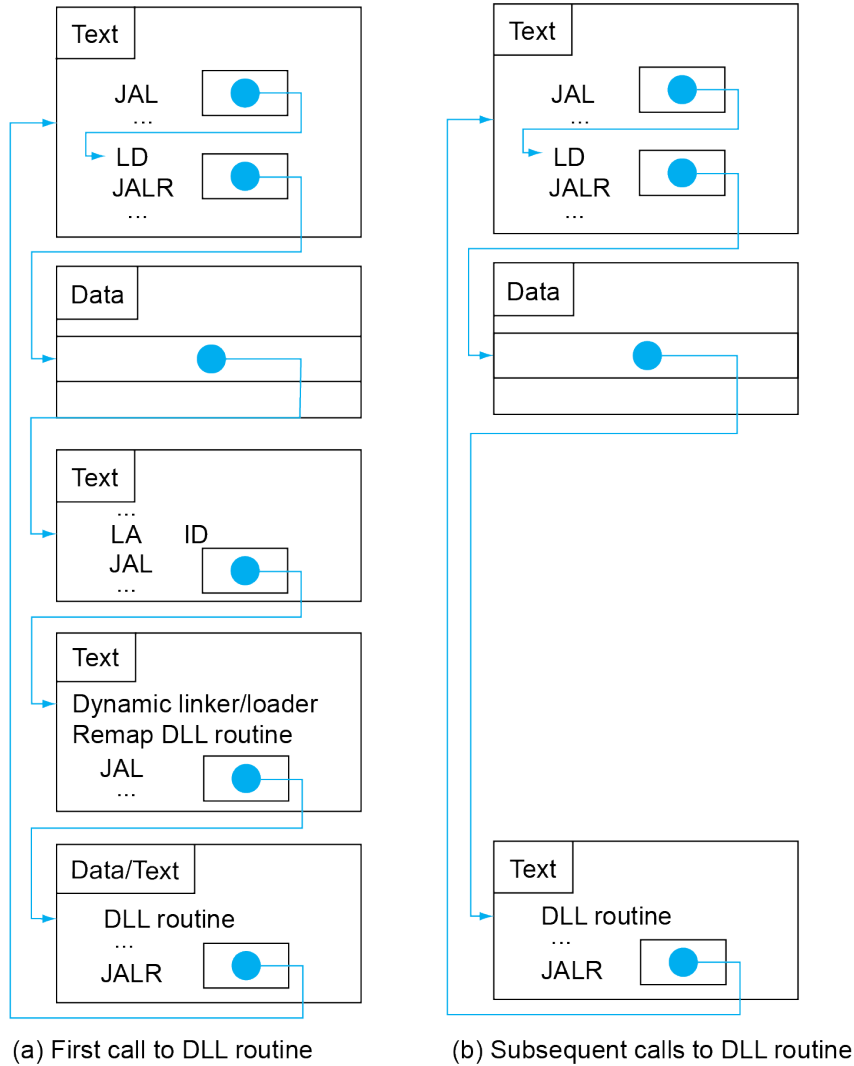
Dynamically  
mapped code



(a) First call to DLL routine



(b) Subsequent calls to DLL routine



Text

```
jal printf();
printf ()
{ ID   x5, printf_addr
  JALR x0, 0(x5)}
```

Data

printf\_addr.word L1

Data

printf\_addr 0x400000

Text

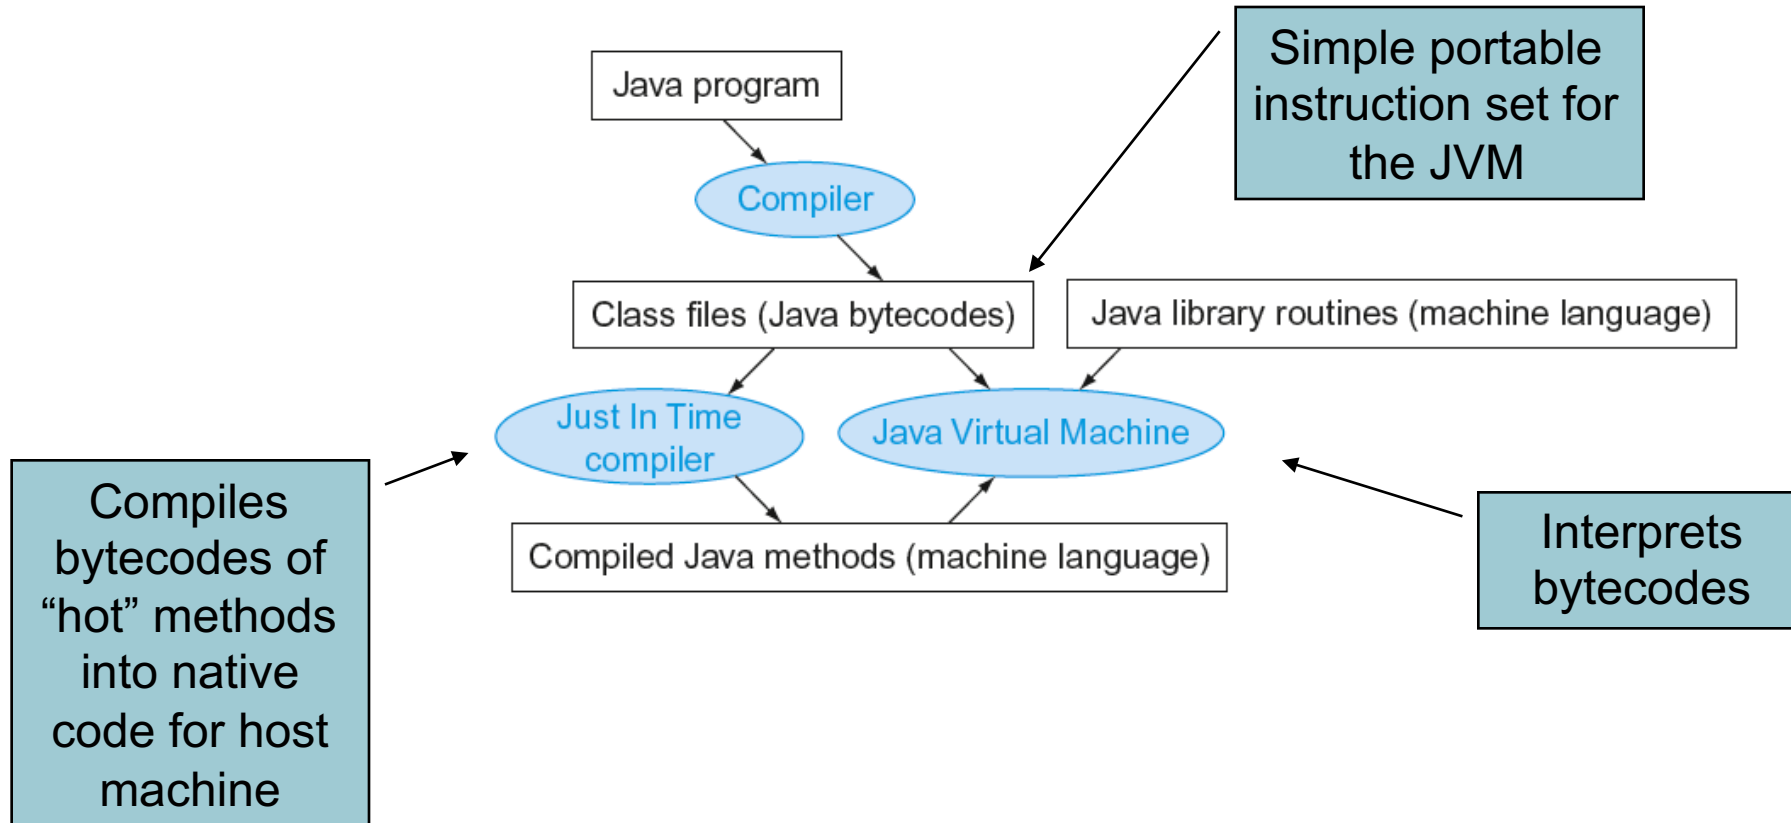
```
L1: LA ID
     JAL x1, DLL;
```

Text

```
printf()
.....
JALR x0, 0(x1)
```

Assume this is loaded into  
0x400000

# Starting Java Applications



# MIPS Instructions

- MIPS: commercial predecessor to RISC-V
- Similar basic set of instructions
  - 32-bit instructions
  - 32 general purpose registers, register 0 is always 0
  - 32 floating-point registers
  - Memory accessed only by load/store instructions
    - Consistent use of addressing modes for all data sizes
- Different conditional branches
  - For <, <=, >, >=
  - RISC-V: blt, bge, bltu, bgeu
  - MIPS: slt, sltu (set less than, result is 0 or 1)
    - Then use beq, bne to complete the branch

# Instruction Encoding

## Register-register

	31	25	24	20	19	15	14	12	11	7	6	0
RISC-V	funct7(7)				rs2(5)		rs1(5)		funct3(3)	rd(5)		opcode(7)
	31	26	25	21	20	16	15	11	10	6	5	0
MIPS	Op(6)			Rs1(5)		Rs2(5)		Rd(5)		Const(5)		Opx(6)

## Load

	31	20	19	15	14	12	11	7	6	0		
RISC-V	immediate(12)				rs1(5)		funct3(3)	rd(5)		opcode(7)		
	31	26	25	21	20	16	15				0	
MIPS	Op(6)			Rs1(5)		Rs2(5)		Const(16)				

## Store

	31	25	24	20	19	15	14	12	11	7	6	0
RISC-V	immediate(7)				rs2(5)		rs1(5)		funct3(3)	immediate(5)		opcode(7)
	31	26	25	21	20	16	15					0
MIPS	Op(6)			Rs1(5)		Rs2(5)		Const(16)				

## Branch

	31	25	24	20	19	15	14	12	11	7	6	0
RISC-V	immediate(7)				rs2(5)		rs1(5)		funct3(3)	immediate(5)		opcode(7)
	31	26	25	21	20	16	15					0
MIPS	Op(6)			Rs1(5)		Opx/Rs2(5)		Const(16)				

# The Intel x86 ISA

- Evolution with backward compatibility
  - 8080 (1974): 8-bit microprocessor
    - Accumulator, plus 3 index-register pairs
  - 8086 (1978): 16-bit extension to 8080
    - Complex instruction set (CISC)
  - 8087 (1980): floating-point coprocessor
    - Adds FP instructions and register stack
  - 80286 (1982): 24-bit addresses, MMU
    - Segmented memory mapping and protection
  - 80386 (1985): 32-bit extension (now IA-32)
    - Additional addressing modes and operations
    - Paged memory mapping as well as segments



# The Intel x86 ISA

- Further evolution...
  - i486 (1989): pipelined, on-chip caches and FPU
    - Compatible competitors: AMD, Cyrix, ...
  - Pentium (1993): superscalar, 64-bit datapath
    - Later versions added **MMX** (Multi-Media eXtension) instructions ~ 57 instructions
    - The infamous FDIV bug
  - Pentium Pro (1995), Pentium II (1997)
    - New microarchitecture (see Colwell, *The Pentium Chronicles*)
  - Pentium III (1999)
    - Added **SSE** (Streaming SIMD Extensions) and associated registers ~ **70 instructions**
  - Pentium 4 (2001)
    - New microarchitecture
    - Added **SSE2** instructions ~ **144 instructions**

# The Intel x86 ISA

- And further...
  - AMD64 (2003): extended architecture to 64 bits
  - EM64T – Extended Memory 64 Technology (2004)
    - AMD64 adopted by Intel (with refinements)
    - Added SSE3 instructions
  - Intel Core (2006)
    - Added SSE4 instructions, virtual machine support
  - AMD64 (announced 2007): SSE5 instructions
    - Intel declined to follow, instead...
  - Advanced Vector Extension (announced 2008)
    - Longer SSE registers, more instructions
- If Intel didn't extend with compatibility, its competitors would!
  - Technical elegance ≠ market success

# Basic x86 Registers

Name	31	0	Use
EAX			GPR 0
ECX			GPR 1
EDX			GPR 2
EBX			GPR 3
ESP			GPR 4
EBP			GPR 5
ESI			GPR 6
EDI			GPR 7
	CS		Code segment pointer
	SS		Stack segment pointer (top of stack)
	DS		Data segment pointer 0
	ES		Data segment pointer 1
	FS		Data segment pointer 2
	GS		Data segment pointer 3
EIP			Instruction pointer (PC)
EFLAGS			Condition codes

# Basic x86 Addressing Modes


- Two operands per instruction

Source/dest operand	Second source operand
Register	Register
Register	Immediate
Register	Memory
Memory	Register
Memory	Immediate

- Memory addressing modes
  - Address in register
  - $\text{Address} = R_{\text{base}} + \text{displacement}$
  - $\text{Address} = R_{\text{base}} + 2^{\text{scale}} \times R_{\text{index}}$  (scale = 0, 1, 2, or 3)
  - $\text{Address} = R_{\text{base}} + 2^{\text{scale}} \times R_{\text{index}} + \text{displacement}$

# x86 Instruction Encoding


a. JE EIP + displacement




b. CALL




c. MOV EBX, [EDI + 45]



d. PUSH ESI



e. ADD EAX, #6765



f. TEST EDX, #42



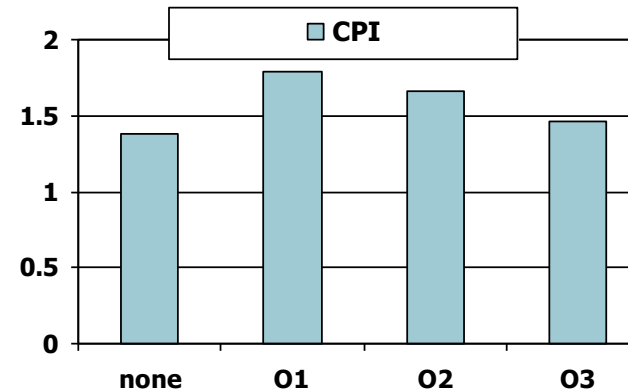



- Variable length encoding
  - Postfix bytes specify addressing mode
  - Prefix bytes modify operation
    - Operand length, repetition, locking, ...

# Implementing IA-32




- Complex instruction set makes implementation difficult
  - Hardware translates instructions to simpler microoperations
    - Simple instructions: 1–1
    - Complex instructions: 1–many
  - Microengine similar to RISC
  - Market share makes this economically viable
- Comparable performance to RISC
  - Compilers avoid complex instructions

# Effect of Compiler Optimization

Compiled with gcc for Pentium 4 under Linux



# Effect of Language and Algorithm






# Fallacies

- Powerful instruction  $\Rightarrow$  higher performance
  - Fewer instructions required
  - But complex instructions are hard to implement
    - May slow down all instructions, including simple ones
  - Compilers are good at making fast code from simple instructions
- Use assembly code for high performance
  - But modern compilers are better at dealing with modern processors
  - More lines of code  $\Rightarrow$  more errors and less productivity

# Fallacies

- Backward compatibility  $\Rightarrow$  instruction set doesn't change
  - But they do accrete more instructions



x86 instruction set